

# House Prediction

Team

2024-12-01

## R Markdown

```
# Load libraries  
library(tidyverse)
```

```
## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --  
## v dplyr      1.1.4      v readr      2.1.5  
## v forcats    1.0.0      v stringr    1.5.1  
## v ggplot2    3.5.1      v tibble     3.2.1  
## v lubridate  1.9.3      v tidyr      1.3.1  
## v purrr      1.0.2  
## -- Conflicts ----- tidyverse_conflicts() --  
## x dplyr::filter() masks stats::filter()  
## x dplyr::lag()     masks stats::lag()  
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(caret)
```

```
##      lattice  
##  
##      'caret'  
##  
## The following object is masked from 'package:purrr':  
##  
##      lift
```

```
library(gbm)
```

```
## Loaded gbm 2.2.2  
## This version of gbm is no longer under development. Consider transitioning to gbm3, https://github.com
```

```
# Read and prepare data  
housing_data <- read.csv("C:/Users/Bowen/Downloads/437/CPTS_437_data/whitman_property_details.csv")  
  
# Enhanced data cleaning with more features  
clean_data <- housing_data %>%  
  mutate(  
    Total_Area = as.numeric(gsub(",", "", ifelse(Total_Area == "None", NA, Total_Area))),
```

```

Year_Built = as.numeric(ifelse(Year_Built == "None", NA, Year_Built)),
Total_Value = as.numeric(gsub(",", "", ifelse(Total_Value == "None", NA, Total_Value))),
Bedrooms = as.numeric(ifelse(Bedrooms == "None", NA, Bedrooms)),
Bathrooms = as.numeric(ifelse(Bathrooms == "None", NA, Bathrooms)),
Garage_Stalls = as.numeric(ifelse(Garage_Stalls %in% c("None", "Block"), 0, Garage_Stalls))
) %>%
filter(!is.na(Total_Value) & !is.na(Total_Area) & !is.na(Year_Built))

## Warning: There were 4 warnings in `mutate()`.
## The first warning was:
## i In argument: `Total_Area = as.numeric(gsub(",", "", ifelse(Total_Area ==
##   "None", NA, Total_Area)))`.
## Caused by warning:
## !      NA
## i Run `dplyr::last_dplyr_warnings()` to see the 3 remaining warnings.

# Enhanced feature engineering with additional features
model_data <- clean_data %>%
  mutate(
    log_value = log(Total_Value + 1),
    log_area = log(Total_Area + 1),
    age = 2024 - Year_Built,
    age_squared = age^2,
    has_garage = ifelse(is.na(Garage_Stalls), 0, 1),
    garage_value = ifelse(is.na(Garage_Stalls), 0, Garage_Stalls),
    condition_score = case_when(
      grepl("3.0", Condition) ~ 3.0,
      grepl("3.5", Condition) ~ 3.5,
      grepl("4.0", Condition) ~ 4.0,
      TRUE ~ 3.0
    ),
    bedrooms = ifelse(is.na(Bedrooms), median(Bedrooms, na.rm=TRUE), Bedrooms),
    bathrooms = ifelse(is.na(Bathrooms), median(Bathrooms, na.rm=TRUE), Bathrooms),
    rooms_per_area = (bedrooms + bathrooms) / log_area,
    age_condition_interaction = age * condition_score
  ) %>%
  select(log_value, log_area, age, age_squared, has_garage, garage_value,
         bedrooms, bathrooms, condition_score, rooms_per_area,
         age_condition_interaction) %>%
  na.omit()

# Split data with stratification
set.seed(123)
train_index <- createDataPartition(model_data$log_value, p = 0.8, list = FALSE)
train_data <- model_data[train_index, ]
test_data <- model_data[-train_index, ]

# Train GBM model with tuned parameters
gbm_model <- gbm(
  log_value ~ .,
  data = train_data,
  distribution = "gaussian",
  n.trees = 1000,


```

```

interaction.depth = 4,
shrinkage = 0.01,
n.minobsinnode = 10,
bag.fraction = 0.8, # Added bagging
cv.folds = 5        # Added cross-validation
)

# Find optimal number of trees
best_iter <- gbm.perf(gbm_model, method = "cv")

```



```

print(paste("Optimal number of trees:", best_iter))

```

```
## [1] "Optimal number of trees: 973"
```

```

# Make predictions using optimal number of trees
predictions <- exp(predict(gbm_model, test_data, n.trees = best_iter)) - 1
actual_values <- exp(test_data$log_value) - 1

# Calculate metrics
rmse <- sqrt(mean((predictions - actual_values)^2))
r2 <- 1 - sum((actual_values - predictions)^2) /
      sum((actual_values - mean(actual_values))^2)
mae <- mean(abs(predictions - actual_values))

```

```
# Print comprehensive metrics
print(paste("RMSE:", format(rmse, scientific = FALSE)))
```

```
## [1] "RMSE: 372187.9"
```


```
print(paste("R-squared:", round(r2, 3)))
```

```
## [1] "R-squared: 0.631"
```

```
print(paste("MAE:", format(mae, scientific = FALSE)))
```

```
## [1] "MAE: 107127.8"
```

```
# Enhanced visualization
importance <- summary(gbm_model, n.trees = best_iter)
```




```
importance_df <- data.frame(
  Feature = importance$var,
  Importance = importance$rel.inf
)

ggplot(importance_df, aes(x = reorder(Feature, Importance), y = Importance)) +
```

```

geom_bar(stat = "identity", fill = "steelblue") +
coord_flip() +
theme_minimal() +
theme(
  plot.title = element_text(face = "bold"),
  axis.text = element_text(size = 10),
  axis.title = element_text(size = 12)
) +
labs(
  x = "Features",
  y = "Relative Importance (%)",
  title = "Feature Importance in House Price Prediction",
  subtitle = paste("Model R2 =", round(r2, 3)),
  caption = paste("RMSE: $", format(rmse, big.mark=",", scientific=FALSE))
)

```



```

# Add residual plot
residuals <- actual_values - predictions
ggplot(data.frame(fitted = predictions, residuals = residuals),
  aes(x = fitted, y = residuals)) +
  geom_point(alpha = 0.5) +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  theme_minimal() +
  labs(
    x = "Fitted Values",

```

```
y = "Residuals",  
title = "Residual Plot",  
)
```

